

令和5年度 修士論文

# Traffic Intersection Negotiation Using Multi-Agent Policy Optimization

Main Supervisor

Tatsuya Yamazaki

Niigata University

Graduate School of Science and Technology

Electrical and Information Engineering

Information Engineering

KHEANG Kim Ang

# Abstract

This study investigates negotiation tasks at intersections for autonomous vehicles and proposes a MARL approach to tackle this challenge.

Negotiating interactions at intersections is crucial for safe and efficient autonomous driving, requiring intelligent decision-making. However, traditional approaches, such as Proximal Policy Optimization (PPO) with Gaussian distribution, face limitations in capturing the complex dynamics of negotiation scenarios. We introduce the utilization of the Beta distribution as an alternative to the Gaussian distribution. The Beta distribution offers several advantages, including greater flexibility in modeling continuous actions and a more accurate representation of the uncertainty associated with negotiation tasks. By leveraging the properties of the Beta distribution, we can enhance the agents' decision-making capabilities, leading to improved negotiation outcomes.

Our approach employs a centralized learning phase and decentralized execution. In the centralized learning phase, agents learn from a global perspective, considering the interactions and dynamics of all vehicles at the intersection. This enables the agents to learn effective negotiation strategies and understand the impact of their actions on the overall traffic flow. During decentralized execution, agents make decisions independently in real time, allowing for adaptive negotiation behaviors based on local observations.

Moreover, we focus on a dual-channel architecture that predicts both the alpha and beta values of the Beta distribution. This architecture provides fine-grained control over the agents' behavior, allowing them to dynamically adjust their negotiation strategies based on the specific context of the intersection and the behavior of other vehicles. By accurately predicting the alpha and beta values, the agents can make informed decisions and exhibit more nuanced negotiation behaviors.

To further improve negotiation performance, we introduce a shared value advantage function that incorporates the joined actions and states of the agents. This shared value function captures the collective impact of all agents, promoting coordination and cooperation in negotiation tasks. By considering the holistic effects of the agents' actions, our approach leads to better negotiation outcomes and improved traffic flow efficiency.

The experiment is carried out through simulation which is a low-cost option to train autonomous vehicles. Through extensive experimentation and comparisons, we demonstrate that our proposed approach outperforms PPO with Gaussian distribution in terms of

learning rate and overall performance. The utilization of the Beta distribution, combined with the dual-channel architecture and the shared value advantage function, enables our agents to exhibit interesting emergent behaviors at intersections. These behaviors include slowing down to let other vehicles pass or strategically cutting through intersections, showcasing the agents' adaptability and effective negotiation skills.

Overall, this thesis presents a comprehensive MARL approach for negotiation tasks at intersections for autonomous vehicles. By utilizing the Beta distribution, employing a centralized learning decentralized execution framework, adopting a dual-channel architecture, and incorporating a shared value advantage function, we achieve superior negotiation performance. The emergent behaviors exhibited by the agents highlight the efficacy and adaptability of our approach in complex negotiation scenarios.

Our findings contribute to the advancement of negotiation tasks in intersection scenarios for autonomous vehicles, facilitating safer and more efficient traffic flow. By leveraging multi-agent reinforcement learning techniques and addressing the limitations of traditional approaches, we pave the way for the development of intelligent and socially aware autonomous vehicles capable of effectively navigating intersections and engaging in successful negotiation protocols.

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Background . . . . .	3
1.2	Motivation . . . . .	4
1.3	Problem Statement . . . . .	5
1.4	Objectives and Scope of the Study . . . . .	5
<b>2</b>	<b>Literature Review</b>	<b>8</b>
2.1	Reinforcement Learning in Traffic Intersection . . . . .	8
2.2	Reinforcement Learning Algorithms . . . . .	10
2.2.1	Model-Free Reinforcement Learning . . . . .	10
2.2.2	Model-Based Reinforcement Learning . . . . .	15
2.3	Single Agent Reinforcement Learning . . . . .	20
2.4	Multi-Agent Reinforcement Learning . . . . .	21
2.5	Data and Simulator Choices . . . . .	22
2.5.1	Gym-Graph-Traffic . . . . .	22
2.5.2	Highway-env . . . . .	24
2.5.3	MetaDrive . . . . .	25
<b>3</b>	<b>Methodology</b>	<b>29</b>
3.1	Markov Decision Process . . . . .	29
3.1.1	Elements of Dec-POMDPs: . . . . .	30
3.2	Gaussian Distribution . . . . .	31
3.3	Policy Gradient Method with Gaussian Approximation: . . . . .	31
3.3.1	Parameterized Policy . . . . .	32

3.3.2	Prediction of Mean Using Neural Network . . . . .	32
3.3.3	Adjusting Standard Deviation Over Time . . . . .	32
3.4	Beta Distribution . . . . .	32
3.5	Policy Gradient Method with Beta Distribution: . . . . .	33
3.5.1	Predicting Alpha and Beta Values from Neural Network: . . . . .	33
3.5.2	Interpretation of $1 - \alpha$ and $1 - \beta$ : . . . . .	34
3.6	Beta Policy . . . . .	34
3.7	Advantage Function (Experience Sharing) . . . . .	35
3.8	Objective Function . . . . .	36
3.9	Centralized Training and Decentralized Execution . . . . .	37
<b>4</b>	<b>Result and Analysis</b>	<b>38</b>
4.1	Experimental Setup . . . . .	38
4.2	Reward Shaping . . . . .	40
4.3	Single Agent Result . . . . .	41
4.4	Multi Agent Result . . . . .	43
4.4.1	Multi-agent PPO Beta . . . . .	44
4.4.2	Multi-agent PPO with Gaussian Distribution . . . . .	44
<b>5</b>	<b>Conclusion</b>	<b>45</b>

# List of Figures

2.1	Gym-graph-traffic Simulator . . . . .	23
2.2	highway-env Simulator . . . . .	25
2.3	MetaDrive Simulator . . . . .	26
3.1	Normal Distribution . . . . .	31
3.2	Beta Distribution . . . . .	33
3.3	Beta Policy . . . . .	34
3.4	$Q$ value Approximator (Critics) . . . . .	35
3.5	Centralized Training and Decentralized Execution . . . . .	37
4.1	Comparison of single agent training . . . . .	41
4.2	Comparison of single agent training by average success rate . . . . .	42
4.3	Comparison of single agent training by average safety rate . . . . .	42
4.4	Comparison of multi-agent training by average reward . . . . .	43

# List of Tables

2.1	Comparison of 'gym-graph-traffic,' 'Highway-env,' and 'MetaDrive' for MARL Research . . . . .	27
4.1	Experiment Setup with PC Specifications . . . . .	38
4.2	PPO Configuration . . . . .	39

# Chapter 1

## Introduction

### 1.1 Background

Reinforcement learning (RL) has emerged as a transformative paradigm in the field of control engineering, enabling agents to learn how to act in complex environments through observation and interaction. This learning approach has seen significant advancements in recent years, largely driven by the exponential growth in computing capability. The integration of deep learning with RL has further accelerated progress, allowing agents to be trained end-to-end by directly processing raw sensor data through neural networks to predict appropriate actions [1]–[5].

Within the realm of RL, various algorithms have been developed [4], [6]–[8], each dictating how an agent learns from its environment. These algorithms can broadly be categorized as model-based and model-free methods. In the context of autonomous vehicles (AVs) navigating intersections, agents must acquire the skills to negotiate interactions effectively. They should be capable of adapting their behavior to be either collaborative, cooperative, or competitive, depending on the situation, in order to safely and efficiently cross intersections.

However, a notable limitation of conventional RL algorithms is that they were primarily designed with the assumption of single-agent environments, where one agent interacts with the environment in isolation. In contrast, intersection scenarios involve multiple agents (vehicles) interacting with one another, necessitating the development of multi-agent reinforcement learning (MARL) techniques [9].



Efforts have been made to extend model-free methods for MARL purposes, aiming to enable agents to learn in collaborative and competitive multi-agent environments. Nevertheless, many of these MARL methods have tended to simplify policy architectures, overlooking essential aspects such as the choice of probability distribution density function or the design of agents' policy architectures.

## 1.2 Motivation

The motivation behind this research stems from the urgent need to address the challenges associated with negotiation tasks at intersections for AVs. As autonomous driving technology continues to advance rapidly, intersections remain one of the most complex and critical environments for AVs to navigate safely and efficiently. Traditional approaches, built on single-agent assumptions and Gaussian distributions, fall short in capturing the intricacies of multi-agent negotiation scenarios.

The potential benefits of effective intersection negotiation are significant, including enhanced traffic flow, reduced congestion, and increased road safety. By enabling AVs to negotiate intersections collaboratively, competitively, or in mixed settings, we can create a more harmonious and efficient traffic ecosystem. Furthermore, successful negotiation protocols at intersections are crucial for the widespread adoption and acceptance of autonomous driving technology.

The adoption of MARL and the utilization of the Beta distribution present promising avenues to revolutionize intersection negotiation. The ability to fine-tune negotiation strategies, consider the joint impact of all agents, and accurately model uncertainty using the Beta distribution are all factors that can greatly improve the negotiation capabilities of AVs.

The research seeks to push the boundaries of existing techniques and bridge the gap between single-agent and MARL. By designing a robust and adaptable negotiation framework, we aim to empower AVs to make well-informed and context-aware decisions at intersections, leading to safer, smoother, and more efficient traffic flow.

Ultimately, this study's motivation lies in contributing to the development of socially-aware and intelligent AVs that can effectively interact with other agents and adapt to various negotiation scenarios at intersections. By overcoming the limitations of conventional

approaches, we aspire to pave the way for the widespread integration of AVs in urban environments, revolutionizing transportation and advancing the future of mobility.

### 1.3 Problem Statement

The problem of negotiating interactions at intersections for AVs poses significant challenges in the field of autonomous driving. Traditional RL algorithms, primarily designed for single-agent environments, struggle to address the complexities of multi-agent negotiation scenarios. While some efforts have been made to extend model-free methods for MARL purposes, these approaches often overlook crucial aspects, such as the choice of probability distribution functions and the design of policy architectures. Additionally, conventional RL algorithms typically rely on Gaussian distributions, which may not adequately capture the complexities of negotiation tasks. As a result, there is a need for a novel and sophisticated approach that leverages the advantages of MARL, centralized learning with decentralized execution, and the Beta distribution to effectively train autonomous agents to negotiate intersections collaboratively, competitively, or in a mixed setting. Addressing these challenges will significantly advance the capabilities of AVs, leading to safer and more efficient traffic flow at intersections.

### 1.4 Objectives and Scope of the Study

The primary objective of this study is to develop and evaluate a novel MARL approach for negotiating interactions at intersections for AVs. The specific objectives are as follows:

- To design a robust and efficient MARL framework that enables AVs to navigate intersections safely and efficiently by learning collaborative, competitive, and mixed negotiation strategies.
- To explore the advantages of the Beta distribution over traditional Gaussian distributions in modeling continuous actions and uncertainty during negotiation tasks, thereby enhancing the adaptability and decision-making capabilities of the autonomous agents.
- To propose a dual-channel architecture for predicting both the  $\alpha$  and  $\beta$  values of the Beta distribution, empowering the agents to fine-tune their negotiation strategies based on the specific context of the intersection and the behavior of other vehicles.

- To develop a shared value advantage function that incorporates joint actions and states of the agents, promoting coordination and cooperation among the AVs during negotiation scenarios, leading to improved negotiation outcomes and overall traffic flow efficiency.
- To conduct extensive experiments and comparisons with conventional PPO algorithms using Gaussian distributions to demonstrate the superior performance and learning rate achieved by our proposed methodology.

This work focuses on model-free unsignalized intersections with no sign where only AVs exist in a homogenous environment. The agents learn together in a mixed cooperative and competitive nature to ensure that the intersection crossing is safe and efficient. This is known as decentralized learning as each agent has its own observation and make decision separately which can ensure emergent behaviors like in the case of multi-agent hide and seek and more [10]–[13]. However, state-of-the-art MARL methods often overlook the importance of agent policy architecture if it is suitable for environment and try to make the learning generalized. This work focuses on extending the policy gradient method to multi-agent version and using Beta distribution instead of standard Gaussian distribution instead. We also put emphasis on agent’s policy architecture designed specifically for Beta distribution. The architecture predicts  $\alpha, \beta$  values to feed to distribution density function in order to sample steering angle and acceleration value to control the vehicle. The research scope encompasses the development of an MARL approach, specifically centered on Proximal Policy Optimization (PPO) with continuous actions and the utilization of the Beta distribution. The investigation will primarily involve AVs acting as the agents, aiming to negotiate interactions with other vehicles in a dynamic traffic environment. The scope includes the exploration of various policy architectures, with particular emphasis on the dual-channel architecture for predicting the  $\alpha$  and  $\beta$  values of the Beta distribution. The shared value advantage function will also be integrated to enable agents to consider the collective impact of their actions and enhance coordination during negotiation. The performance of the proposed methodology will be evaluated through extensive simulations and comparisons with traditional PPO algorithms utilizing Gaussian distributions. The evaluation metrics will include learning rate, negotiation outcomes, and the emergence of interesting behaviors such as slowing down to let other vehicles pass or cutting through

intersections. It is important to note that this study focuses specifically on the intersection negotiation tasks for AVs and does not cover other aspects of autonomous driving, such as motion planning or perception. Additionally, while the study aims to achieve superior performance in negotiation scenarios, it may not address all possible scenarios that AVs may encounter in real-world urban environments.

## Chapter 2

# Literature Review

### 2.1 Reinforcement Learning in Traffic Intersection

RL has shown great promise in enabling vehicles to make intelligent decisions on the road. RL algorithms allow AVs to learn from their interactions with the environment, such as other vehicles, pedestrians, traffic lights, and road conditions. This learning process helps the vehicles adapt to changing situations and make appropriate decisions in real-time.

One of the significant advantages of RL in autonomous driving is its ability to handle complex and uncertain environments. RL agents can navigate through various traffic scenarios, including intersections, merging lanes, and lane-changing maneuvers [14], [15], by learning optimal policies based on trial and error. RL also allows AVs to adapt to diverse driving conditions, such as different weather conditions and traffic patterns.

Applications of RL in autonomous driving include:

- **Trajectory Planning:** RL can be used to plan safe and efficient trajectories for AVs, robots [16], [17], considering various factors such as traffic density, speed limits, and road conditions.
- **Behavior Prediction:** RL algorithms can predict the behavior of other road users, such as pedestrians [18] and human-driven vehicles, to anticipate their actions and plan accordingly.
- **Lane Keeping and Lane Changing:** RL can help AVs maintain their lane position

and make safe lane-changing maneuvers [14] based on surrounding traffic.

- **Intersection Negotiation:** As seen in the context of this study, RL is applied to enable AVs to negotiate intersections [19]–[21] effectively, making decisions on when to yield or proceed.
- **Optimizing Energy Consumption Efficiency:** RL can be used to optimize the energy consumption [22], [23] of AVs, leading to longer battery life and reduced carbon emissions.
- **Traffic Flow Optimization:** RL can be employed to optimize traffic flow [24] and reduce congestion in urban environments, improving overall transportation efficiency.
- **Autonomous Parking:** RL can be used to enable AVs to park themselves [25] in parking lots or on the street.

The use of RL in autonomous driving continues to evolve rapidly, with ongoing research and development focusing on improving safety, efficiency, and adaptability. As autonomous driving technology matures, RL is expected to play a crucial role in enhancing the capabilities of AVs and making them more reliable and intelligent on the road. Many works try to improve transportation safety and efficiency at the intersection by integrating various RL techniques such as intelligent traffic light (sometimes known as adaptive signal control) [26]–[28]. However, as AVs are gaining ground, problems arose as AVs had difficulty trying to recognize the state of traffic signal in occluded situation [29]. AVs can also rely on other data such as Lidar points, and other sensors to navigate through the intersection which renders the adaptive signal control unsuitable for AVs. Another approach was to place an intersection controller (IC) somewhere at the center of the intersection. The IC relies on various data such as video footage and lane detectors to coordinate vehicle movement. AVs need to request passage to the IC in a similar fashion to air traffic control [19], [30]–[32]. As the solution becomes more complex, the infrastructure cost also increased which makes the situation only available in a big city. RL brings promising solutions to solve control problems in this ever-growing big-data community.

## 2.2 Reinforcement Learning Algorithms

Typically, RL algorithms are divided into 2 schools of thought —model-based and model-free [33] approaches. Each class has several sub-branches that address specific challenges and requirements in different problem domains.

### 2.2.1 Model-Free Reinforcement Learning

Most world problems evolve round uncertainty where agents often have incomplete view of the environment. This is where model-free RL emerges. Model-free RL is a type of RL approach where an agent learns to make decisions and take actions in an environment without explicitly building a model of the environment. In traditional RL, an agent typically constructs a model that represents the transition dynamics and rewards of the environment. This model helps the agent to simulate and plan future actions to make decisions. However, in model-free RL, the agent directly learns from experience through trial and error without explicitly constructing a model.

In model-free RL, the agent interacts with the environment over multiple episodes or time steps, and it receives feedback in the form of rewards based on the actions it takes. The primary goal of the agent is to maximize the total accumulated reward over time. To achieve this, the agent uses different algorithms and techniques to learn the best policy, which is a mapping from states to actions that guides the agent in making decisions.

Model-free RL is particularly useful when the environment is complex, and constructing an accurate model is difficult or computationally expensive. Instead of relying on a model, model-free RL agents learn from real interactions with the environment, making them more adaptable and applicable to a broader range of scenarios. However, model-free RL also has some challenges. Since it does not use a model, learning can be more sample inefficient [33], requiring a large number of interactions with the environment to converge to an optimal policy. Additionally, model-free RL can sometimes struggle with environments that have sparse rewards or deceptive reward structures. Nevertheless, with proper tuning and algorithm selection, model-free RL can be highly effective in a wide range of real-world applications.

## Q-Learning

Q-learning is a model-free RL algorithm that estimates the action-value function  $Q(s, a)$ , representing the expected cumulative reward when taking action  $a$  in state  $s$ . It is a widely used algorithm for solving Markov Decision Processes (MDPs) without requiring a model of the environment's dynamics.

The Q-value is updated iteratively using the Bellman equation:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

where:  $Q(s, a)$  is the Q-value of state-action pair  $(s, a)$ ,  $r$  is the immediate reward received after taking action  $a$  in state  $s$ ,  $\alpha$  is the learning rate that controls the step size of updates,  $\gamma$  is the discount factor that balances future rewards' importance,  $s'$  is the next state reached after taking action  $a$  in state  $s$ , and  $\max_{a'} Q(s', a')$  represents the maximum Q-value over all possible actions in state  $s'$ .

The Q-learning algorithm follows the following steps:

1. Initialize the Q-values for all state-action pairs to arbitrary values.
2. Observe the current state  $s$ .
3. Select an action  $a$  using an exploration policy, such as  $\epsilon$ -greedy, which balances between exploration and exploitation.
4. Execute the selected action  $a$  in the environment.
5. Observe the immediate reward  $r$  and the next state  $s'$ .
6. Update the Q-value for the current state-action pair using the Bellman equation.
7. Repeat steps 2 to 6 until convergence or a predefined number of iterations.

Q-learning is an off-policy algorithm, meaning that it learns the optimal policy while following a different behavior policy for exploration. The exploration policy allows the agent to explore the environment and learn about different state-action pairs. As Q-learning iteratively updates the Q-values, it converges to the optimal action-value function  $Q^*(s, a)$ , which represents the maximum expected cumulative reward from each state-action pair under the optimal policy. Q-learning is known to be model-free, which means



it does not require knowledge of the environment’s dynamics or transition probabilities. Instead, it learns from interacting with the environment and experiencing different state-action-reward sequences. The Q-learning algorithm is widely used due to its simplicity and ability to handle large state spaces. However, it may suffer from slow convergence in complex environments, and variants like Double Q-learning and Deep Q-Networks (DQNs) have been introduced to address these limitations.

## DQNs

DQNs is a deep learning extension of the Q-learning algorithm. DQNs leverage the power of deep neural networks to approximate the action-value function  $Q(s, a)$  in complex and high-dimensional environments. It was introduced as a breakthrough in combining deep learning and RL, particularly in solving Atari 2600 games [34].

The core idea of DQNs is to use a neural network as a function approximator for the action-value function. The neural network takes the state  $s$  as input and outputs a vector of Q-values, one for each possible action. The action with the highest Q-value is then selected as the agent’s action in the given state.

DQNs aim to solve two key challenges in Q-learning:

- **Curse of Dimensionality:** In traditional tabular Q-learning, the size of the Q-table grows exponentially with the number of states and actions, making it infeasible for large state spaces. DQNs address this issue by using a neural network to generalize across states and actions.
- **High-Dimensional Inputs:** Many real-world environments have high-dimensional state spaces, such as images in computer vision tasks. DQNs can handle these inputs by using convolutional neural networks (CNNs) as the function approximator.

The DQN algorithm follows these main steps:

1. **Experience Replay:** DQNs use an experience replay buffer to store transitions (state, action, reward, next state) experienced by the agent. This buffer helps decorrelate experiences and reduce the correlation between consecutive updates, leading to more stable training.
2. **Target Network:** To stabilize the learning process, DQNs use a separate target network with fixed parameters to generate the target Q-values during the update step. The

target network’s parameters are updated periodically to match the primary Q-network.

**3. Loss Function:** The DQN loss function is the mean squared error between the predicted Q-values and the target Q-values, given by the Bellman equation:

$$Loss = \frac{1}{N} \sum_i (Q(s_i, a_i) - (r_i + \gamma \max_{a'} Q_{\text{target}}(s'_i, a')))^2$$

where  $N$  is the batch size,  $Q_{\text{target}}$  represents the target Q-network, and  $\gamma$  is the discount factor.

The DQN algorithm iteratively samples batches of experiences from the replay buffer, computes the loss, and performs stochastic gradient descent to update the primary Q-network’s parameters.

DQNs have demonstrated remarkable success in various tasks, including playing Atari games and controlling robotic systems. They can handle high-dimensional inputs and effectively approximate the action-value function in complex environments.

However, DQNs also have some challenges, such as the tendency to overestimate Q-values and difficulties in handling continuous action spaces. Several extensions, such as Double DQNs and Dueling DQNs, have been proposed to address these limitations and further improve performance.

## Policy Gradient Methods

Policy Gradient Methods are a class of RL algorithms that directly optimize the policy of an agent to maximize the expected cumulative reward. Unlike value-based methods that estimate the action-value function  $Q(s, a)$ , policy gradient methods aim to learn the policy  $\pi(a|s)$ , which is a mapping from states  $s$  to actions  $a$ .

The policy is typically represented as a parametric function, such as a neural network, with trainable parameters  $\theta$ . The objective of policy gradient methods is to find the optimal parameters  $\theta^*$  that maximize the expected reward:

$$J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \gamma^t r_t \right]$$

where  $J(\theta)$  is the objective function,  $\tau$  is a trajectory (sequence of states, actions, and rewards),  $\pi_\theta$  is the policy with parameters  $\theta$ ,  $r_t$  is the reward at time step  $t$ ,  $\gamma$  is the discount factor, and  $T$  is the maximum time step in the trajectory.

Policy gradient methods employ gradient ascent to update the policy parameters in the direction that increases the expected reward. The policy gradient is computed using the likelihood ratio trick:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \left( \sum_{t'=t}^T \gamma^{t'-t} r_{t'} \right) \right]$$

This gradient can be estimated through sampling by interacting with the environment and collecting trajectories. The policy parameters are then updated using stochastic gradient ascent:

$$\theta_{\text{new}} = \theta_{\text{old}} + \alpha \nabla_{\theta} J(\theta)$$

where  $\alpha$  is the learning rate.

Policy gradient methods offer several advantages, such as handling continuous action spaces and directly optimizing the policy without relying on value functions. They are also well-suited for problems with stochastic policies or environments.

However, policy gradient methods can suffer from high variance in the gradient estimates, which can lead to slow convergence. To address this, variance reduction techniques like baseline subtraction and advantage functions are used to stabilize training.

Additionally, modern policy gradient methods, such as Proximal Policy Optimization (PPO) and Trust Region Policy Optimization (TRPO), introduce mechanisms to constrain the policy updates to ensure more stable and controlled learning.

Overall, policy gradient methods have proven to be effective in a wide range of applications, including robotic control, game playing, and natural language processing, where learning directly from high-dimensional observations is crucial.

## Proximal Policy Optimization (PPO)

Proximal Policy Optimization (PPO) is a state-of-the-art policy gradient algorithm for RL. It addresses the challenges of traditional policy gradient methods, such as high variance and instability, by introducing constraints on policy updates to ensure more stable and controlled learning [7].

PPO operates by iteratively optimizing a surrogate objective function that approximates the policy's performance while avoiding drastic policy changes that could lead to divergence.

The objective function is defined as the clipped surrogate objective:

$$L^{CLIP}(\theta) = \mathbb{E}_t \left[ \min \left( \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)} A^{\pi_{\theta_{\text{old}}}}(s_t, a_t), \text{clip} \left( \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}, 1 - \epsilon, 1 + \epsilon \right) A^{\pi_{\theta_{\text{old}}}}(s_t, a_t) \right) \right]$$

where:  $\theta$  represents the current policy parameters,  $\theta_{\text{old}}$  represents the old policy parameters (used as a reference),  $\pi_\theta(a_t|s_t)$  is the probability of taking action  $a_t$  in state  $s_t$  under the current policy,  $\pi_{\theta_{\text{old}}}(a_t|s_t)$  is the probability of taking action  $a_t$  in state  $s_t$  under the old policy,  $A^{\pi_{\theta_{\text{old}}}}(s_t, a_t)$  is the advantage function, representing the advantage of taking action  $a_t$  in state  $s_t$  under the old policy,  $\epsilon$  is a hyperparameter that controls the magnitude of the policy change.

The clipped surrogate objective encourages small policy updates by limiting the change in the likelihood ratio between the new and old policies. This constraint prevents the policy from changing too rapidly, which contributes to the algorithm’s stability.

The PPO algorithm operates in two phases:

1. **Data Collection:** The agent interacts with the environment to collect trajectories using the current policy  $\pi_\theta$ .
2. **Policy Update:** The policy parameters  $\theta$  are updated using stochastic gradient ascent to maximize the clipped surrogate objective function  $L^{CLIP}(\theta)$ .

PPO’s ability to strike a balance between exploration and exploitation makes it effective in a wide range of tasks. The algorithm has been successfully applied to various domains, including robotic control, game playing, and natural language processing.

PPO also offers a robust and efficient solution for parallelizing training across multiple environments, further accelerating learning in complex environments.

Overall, PPO has become a popular choice for deep RL due to its stable performance, ease of implementation, and excellent results across diverse problem domains.

### 2.2.2 Model-Based Reinforcement Learning

Model-based RL is an approach where an agent learns and utilizes an explicit model of the environment to make decisions and plan actions. Unlike model-free RL, which directly learns from experience without constructing a model, model-based RL involves creating a representation of how the environment behaves, including its transition dynamics and

reward function. This model is then used to simulate possible trajectories and predict the outcomes of different actions.

The typical process of model-based RL can be summarized as follows:

- **Model Learning** : The agent collects data by interacting with the environment and observes the outcomes of its actions. It uses this data to build a model of the environment. The model could be in the form of a transition function, which predicts the next state given the current state and action, and a reward function, which estimates the expected reward for a state-action pair.
- **Planning** : The agent uses its model to simulate possible trajectories and predict the outcomes of different actions. It can perform offline planning to evaluate different sequences of actions and predict the total expected rewards. This process allows the agent to make decisions that are based on the simulations and predictions from the model, rather than relying solely on trial and error in the real environment.
- **Policy Improvement** : After simulating different trajectories and estimating the rewards, the agent can update its policy to improve decision-making. It can use various optimization algorithms to find a policy that maximizes the expected cumulative reward based on the predictions from the model.

Model-based RL has several advantages:

- **Sample Efficiency** : Since the agent can simulate interactions with the environment using its model, it can potentially learn more efficiently and require fewer real interactions to find an optimal policy.
- **Safety** : The agent can use its model to explore the environment safely, avoiding potentially harmful or risky actions by simulating their consequences.
- **Handling Partial Observability** : Model-based RL can be effective in partially observable environments where the agent does not have access to complete information about the state. The model helps the agent predict the underlying state, improving decision-making.

However, model-based RL also has its challenges:

- **Model Accuracy** : The success of model-based RL heavily depends on the accuracy of the learned model. If the model is inaccurate, the agent's decisions and predictions can be suboptimal.
- **Model Complexity** : Constructing an accurate model may be computationally expensive or infeasible in complex environments with large state spaces.
- **Planning Overhead** : Planning using the model can introduce additional computational overhead, which might limit the real-time applicability of the approach in certain scenarios.

## Dynamic Programming

is a powerful optimization technique used in computer science and mathematics to solve problems by breaking them down into smaller subproblems and solving each subproblem only once. The solutions to these subproblems are stored and reused when needed, which avoids redundant calculations and leads to more efficient computations.

The fundamental idea behind dynamic programming is to divide a complex problem into smaller, overlapping subproblems and solve each subproblem just once, storing its solution in a table or array. Then, the solution to the original problem is constructed by combining the solutions of the subproblems.

The process of dynamic programming can be described by the following steps:

- **Define the Recurrence Relation** : Express the solution to a larger problem in terms of the solutions to its subproblems. This recursive relationship defines the optimal value of the larger problem based on the optimal values of its subproblems.
- **Identify Overlapping Subproblems** : Analyze the problem to identify if there are any overlapping subproblems, i.e., if the same subproblems are solved multiple times in the computation. This is the key property that makes dynamic programming effective, as we can avoid redundant computations.
- **Create a Memoization Table** : Use a data structure, such as an array or a dictionary, to store the solutions of the subproblems as they are computed. This is called memoization, and it helps in efficiently retrieving solutions for previously solved subproblems.

- **Bottom-Up or Top-Down Approach** : Dynamic programming can be implemented using either a bottom-up or top-down approach.
  - **Bottom-Up** : In the bottom-up approach, we start solving the smallest subproblems first and gradually build up to the solution of the original problem.
  - **Top-Down** : In the top-down approach, also known as memoization, we start with the original problem and break it down into smaller subproblems.

### Monte Carlo Tree Search (MCTS)

Monte Carlo Tree Search (MCTS) is a popular algorithm used in decision-making and game playing problems, particularly in games with large state spaces and complex branching factors. It is commonly applied in artificial intelligence for games like Go, Chess, and other strategic board games. MCTS combines elements of random sampling (Monte Carlo) with tree-based search to efficiently explore the game tree and make informed decisions.

The MCTS algorithm can be explained in the following steps:

1. **Selection** : Starting from the root of the search tree, the algorithm traverses the tree by selecting child nodes based on a selection policy, often using Upper Confidence Bound (UCB) or variants like UCT (Upper Confidence bounds for Trees). The selection policy balances between exploring new nodes and exploiting existing knowledge to focus on promising branches.
2. **Expansion** : Once a leaf node (a node with unexplored children) is reached, the algorithm expands the tree by adding child nodes representing possible game moves or actions.
3. **Simulation (Rollout)** : From the newly added child node, the algorithm performs a Monte Carlo simulation (also known as a rollout) by playing the game randomly until a terminal state (end of the game) is reached. The result of the simulation is a reward or utility value representing the outcome of the game from that state.
4. **Backpropagation** : The reward obtained from the simulation is backpropagated up the tree, updating the statistics of each visited node, such as the total reward and the number of visits. This information helps in refining the selection policy and improves the decision-making.

5. **Repeated Iteration:** Steps 1 to 4 are repeated for a certain number of iterations or until a predefined computational budget is reached. The more iterations performed, the more refined the search tree becomes, leading to better decisions.

## Model Predictive Control (MPC)

Model Predictive Control (MPC) is an advanced control strategy used in various engineering fields, such as process control, robotics, and automotive systems. MPC is particularly well-suited for systems with complex dynamics, constraints, and uncertainties. It involves using a model of the system's behavior to predict future states and optimize a control sequence over a finite time horizon. The control sequence is applied to the system, and the process is iterated in a receding horizon fashion, continually updating the control actions based on the most recent measurements and predictions.

Here's a detailed explanation of Model Predictive Control:

- **Modeling the System :** The first step in MPC is to create a dynamic model of the system being controlled. This model should accurately represent the system's behavior and dynamics. Depending on the complexity of the system and the available information, the model can be a physics-based model, an empirical model, or even a data-driven model based on machine learning techniques.
- **Prediction Horizon :** MPC works by optimizing a control sequence over a finite prediction horizon, which is a predetermined time interval into the future. The length of the prediction horizon determines how far into the future the control actions are computed. Longer prediction horizons can provide better performance, but they also increase the computational burden.
- **Optimization Objective :** The main objective of MPC is to optimize a cost function over the prediction horizon. The cost function is a measure of how well the system performs, and it typically includes components related to control effort, desired setpoints, and constraints. By minimizing or maximizing the cost function, MPC finds the optimal control sequence that achieves the desired control objectives.
- **Constraints :** MPC can handle both hard and soft constraints on the system's states and inputs. Hard constraints are strict limits that cannot be violated, while



soft constraints are desired limits that can be violated if necessary. MPC can also handle probabilistic constraints, which are constraints that must be satisfied with a certain probability.

- **Receding Horizon Control** : Once the optimal control sequence is computed, only the first control action is applied to the system. The process is then repeated, with the optimization being performed again using the most recent measurements and predictions. This is called receding horizon control, as the optimization horizon keeps receding into the future.
- **Real-Time Implementation** : MPC is typically implemented in real-time using a digital computer. The optimization problem is solved at each time step, and the first control action is applied to the system. The process is repeated at the next time step, and so on. The real-time implementation of MPC requires fast and efficient algorithms to solve the optimization problem.
- **Robustness and Adaptation** : MPC is inherently robust to disturbances and uncertainties, as it continually updates the control actions based on the most recent measurements and predictions. MPC can also be extended to handle model uncertainties and unmodeled dynamics by using adaptive MPC techniques.

MPC is widely used in industrial automation, autonomous systems, power systems, and various other applications where precise control and handling of complex dynamics and constraints are essential. It provides robust and adaptive control, making it a valuable tool in controlling complex systems in real-world scenarios.

## 2.3 Single Agent Reinforcement Learning

Single-agent RL algorithms are designed to solve problems where there is only one agent interacting with the environment. There are advantages and disadvantages to using single-agent RL algorithms in autonomous driving applications:

- **Simplicity**: Single-agent RL is relatively straightforward to implement and understand since there is only one agent interacting with the environment.

- **Model-Based Approaches:** Model-based RL algorithms can work effectively in environments where the agent has a complete understanding of the environment’s dynamics, enabling accurate predictions.
- **Sample Efficiency:** Single-agent RL algorithms can be more sample-efficient compared to MARL algorithms since they do not involve interactions among multiple agents.

However, as the number of agents increases, we have some disadvantages:

- **Scalability:** In complex environments with many agents, single-agent RL algorithms may struggle to scale effectively, especially when the agents’ actions affect each other.
- **Non-Stationary Environments:** Single-agent RL assumes a stationary environment, which may not hold true in dynamic settings with multiple agents.

## 2.4 Multi-Agent Reinforcement Learning

The advantages of MARL in autonomous driving include:

- **Coordination and Cooperation:** MARL allows agents to coordinate and cooperate to achieve common goals or address conflicting objectives, making it suitable for problems with multiple agents interacting.
- **Dynamic Environments:** MARL is well-suited for non-stationary and complex environments, where the behavior of agents affects each other and evolves over time.
- **Flexibility:** MARL algorithms can handle various scenarios, from collaborative tasks to competitive interactions, making them versatile in real-world settings.

However, MARL also has some disadvantages:

- **Curse of Dimensionality:** As the number of agents increases, the state space grows exponentially, leading to challenges in scalability and computation.
- **Communication Overhead:** In some MARL algorithms, communication among agents may be required to coordinate actions, introducing additional overhead and complexity.
- **Partial Observability:** In scenarios where agents have limited visibility of the environment or other agents’ states, coordinating effectively becomes more challenging.

As agents interaction became inevitable, several MARL algorithms were proposed [9], [10], [20], [35]–[39]. However, MARL methods often simplify policy architecture and overlooked small parts of policy update such as the type of probability distribution density function or agents’ policy architecture.

## 2.5 Data and Simulator Choices

In this section, we discuss the data and simulator choices for our experiments.

### 2.5.1 Gym-Graph-Traffic

Gym-graph-traffic is an OpenAI Gym-compatible road traffic simulator that provides a platform for studying and experimenting with various traffic scenarios. To use the simulator, users can easily install it from the provided GitHub repository<sup>1</sup> and run examples using Python scripts like *minimal.py*. The simulator’s configuration is customizable through parameters available in the *params.py* file, allowing users to adjust aspects such as episode duration, step length, and possible actions per intersection. The simulator Fig. 2.1 operates on different time units, such as Experiment, Episode, Step, and Update, each representing specific durations in real-time seconds. Actions in the simulator are represented as integers and can be translated to a base 10 representation. By providing observation vectors for the average number of cars and average speed per road segment, users can monitor and evaluate traffic performance. With the option to choose from several preset road network configurations, the simulator offers flexibility in studying various traffic environments. Although not yet supporting Gym-like rendering, rendering options can still be configured within the *params.py* file. Overall, Gym-Graph-Traffic serves as a valuable tool for researchers and developers interested in traffic management and AV behavior.

Gym-Graph-Traffic may not be suitable for MARL research due to its design and focus on single-agent traffic simulation. The environment is primarily centered around road traffic simulation with a single ego-vehicle navigating through the traffic scenarios. While it provides valuable features for studying the behavior of a single agent in various traffic situations, it lacks native support for handling multiple agents and their interactions.

MARL research requires environments that can model and simulate the complex

---

<sup>1</sup><https://github.com/rltraffic/gym-graph-traffic>




Figure 2.1: Gym-graph-traffic Simulator

General interface of Gym-graph-traffic that allows us to simulate urban environment with a lot of entities

interactions between multiple agents, where the actions of each agent can influence the behavior of others. In contrast, Gym-Graph-Traffic is primarily designed for single-agent learning, and its architecture may not easily accommodate the dynamics and coordination required for multi-agent interactions.

To foster multi-agent research, specialized environments are typically required that explicitly handle multiple agents, allow for communication and coordination between agents, and offer capabilities for joint learning or adversarial settings. Such environments must consider the diverse strategies and behaviors that emerge when multiple intelligent agents interact, which may not be fully addressed in Gym-Graph-Traffic’s single-agent-focused simulation.

Researchers interested in MARL would be better served by using environments explicitly designed for this purpose, which provide the necessary tools and support for studying complex agent interactions and cooperation. While Gym-Graph-Traffic serves as a valuable

tool for single-agent traffic simulation, its limitations in multi-agent scenarios make it less suitable for comprehensive MARL research.

### 2.5.2 Highway-env

Highway-env [40] is a collection of environments designed for autonomous driving and tactical decision-making tasks, maintained by Edouard Leurent. The package includes various environments, each representing different driving scenarios. One such environment is *Highway*, where the ego-vehicle navigates a multilane highway with other vehicles. The agent’s objective is to achieve high speed while avoiding collisions with neighboring vehicles and staying on the right side of the road. A faster variant called *highway-fast-v0* is available for large-scale training with reduced simulation accuracy.

Other environments include *Merge*, where the ego-vehicle must maintain high speed and make room for incoming vehicles merging into the traffic, *Roundabout*, requiring the ego-vehicle to pass a roundabout with lane changes and collision avoidance, *Parking*, a goal-conditioned task involving parking the vehicle in a given space with the appropriate heading, *Intersection*, an intersection negotiation task with dense traffic, and *Racetrack*, a continuous control task involving lane-keeping and obstacle avoidance.

The package<sup>2</sup> provides examples of agents that can solve the highway-env environments, such as the DQN, Deep Deterministic Policy Gradient (DDPG), Value Iteration, and MCTS agents. These agents utilize various RL techniques to perform the tasks efficiently. The agent can interact with the environment by taking actions, receiving observations, rewards, and additional information at each step.

Overall, highway-env offers a comprehensive set of environments for researchers and developers to study and develop autonomous driving algorithms and tactical decision-making strategies.

Highway-env Fig. 2.2 may not be ideal for MARL research due to its primary focus on single-agent driving scenarios. The environment is specifically designed for studying the behavior of a single ego-vehicle navigating through various highway and road scenarios. While it provides a range of driving challenges for a lone agent, it lacks built-in support and mechanisms for handling multiple agents and their interactions.

In MARL research, environments need to accommodate complex interactions between

---

<sup>2</sup><https://github.com/Farama-Foundation/HighwayEnv>




Figure 2.2: highway-env Simulator

Top view of the highway-env simulator, where the ego-vehicle (green) navigates an intersection with other vehicles (blue)

multiple agents, where the actions of each agent can influence the behavior of others. Highway-env is not explicitly designed to handle such multi-agent dynamics, and modifying it to support multiple agents could be non-trivial and may not produce optimal results.

For comprehensive MARL research, specialized environments that explicitly cater to multi-agent scenarios, such as cooperative or competitive driving situations, are more appropriate. These environments should facilitate communication and coordination between agents, support joint learning, and provide tools for studying emergent behaviors arising from multi-agent interactions.

### 2.5.3 MetaDrive

MetaDrive [41] is a driving simulator designed for generalizable RL research<sup>3</sup>. It supports generating diverse driving scenarios with various road maps and traffic settings, making it suitable for studying RL agents' generalization capabilities. The simulator is Fig. 2.3 lightweight, easy to install, and can achieve up to 300 FPS on a standard PC. It

<sup>3</sup><https://github.com/metadriverse/metadrive>

provides realistic physics simulation and multiple sensory inputs, including Lidar, RGB images, top-down semantic maps, and first-person view images. Examples of different driving scenarios are provided, such as Single Agent Environment, Multi-Agent Environment, and Real Environment, along with instructions on how to run them. To build the RL environment, users can code in the Gymnasium format using the provided MetaDriveEnv class.



Figure 2.3: MetaDrive Simulator

First person view of MetaDrive simulator showing mini-map, first person view, and top-down view

MetaDrive stands out as an excellent choice for MARL research as in Table 2.1 due to its provision of rich observations and robust multi-agent support. In the context of MARL, having access to diverse and comprehensive observations is crucial for agents to effectively learn and adapt in complex environments. MetaDrive addresses this need by offering multiple sensory inputs, including Lidar, RGB images, top-down semantic maps, and first-person view (FPV) images. These varied observation modalities enable agents to perceive the environment from different perspectives, promoting a more comprehensive understanding of the driving scenarios.

Table 2.1: Comparison of 'gym-graph-traffic,' 'Highway-env,' and 'MetaDrive' for MARL Research

Feature	<b>gym-graph-traffic</b>	<b>Highway-env</b>	<b>MetaDrive</b>
Support for Multi-Agent RL Research	No	No	Yes
Focus	Single-Agent Traffic	Single-Agent Driving	Single and Multi Agent Driving
Rich Observation Modalities	Limited	Limited	Yes (Lidar, RGB images, top-down semantic map, and FPV images)
Multi-Agent Interactions	No	No	Limited (Possible with modification)
Built for Complex Scenarios	Simple Intersections	Simple Highways	Realistic and Diverse Driving Scenarios

Furthermore, MetaDrive’s multi-agent support is a key strength that sets it apart. In multi-agent environments, interactions and cooperation between agents play a significant role in determining their overall performance and behavior. MetaDrive facilitates the simulation of multiple agents simultaneously, allowing researchers to study cooperative, competitive, and mixed interactions between agents. This capability is invaluable for investigating complex real-world scenarios where multiple AVs, pedestrians, and traffic participants interact dynamically. By combining rich observations with multi-agent support, MetaDrive provides a versatile platform for exploring diverse driving scenarios, such as intersection negotiations, cooperative lane merging, and complex traffic situations.



The ability to model and train agents in such realistic and interactive environments fosters the development of more robust and generalizable MARL algorithms, contributing significantly to the advancement of autonomous driving systems and intelligent transportation technologies. As a result, MetaDrive emerges as a superior option for researchers seeking a comprehensive and effective tool for MARL research in driving and traffic-related domains.

## Chapter 3

# Methodology

### 3.1 Markov Decision Process

A Markov decision process (MDP) is a tuple  $(S, A, P, R, \gamma)$ , where:  $S$  is a set of states.  $A$  is a set of actions.  $P(s'|s, a)$  is the probability of transitioning from state  $s$  to state  $s'$  when action  $a$  is taken.  $R(s, a)$  is the reward received when taking action  $a$  in state  $s$ .  $\gamma \in [0, 1]$  is a discount factor that determines the importance of future rewards. The MDP definition can be interpreted as follows:

- The state  $s$  represents the current state of the environment.
- The action  $a$  represents the agent's decision in the current state.
- The transition probability  $P(s'|s, a)$  represents the probability of the environment transitioning to state  $s'$  after the agent takes action  $a$  in state  $s$ .
- The reward  $R(s, a)$  represents the amount of reward the agent receives after taking action  $a$  in state  $s$ .
- The discount factor  $\gamma$  represents how much the agent values future rewards relative to immediate rewards.

Decentralized Partially Observable Markov Decision Processes (Dec-POMDPs) are a class of MARL problems where multiple agents collectively aim to maximize a global objective function while facing uncertainty in a partially observable environment.

### 3.1.1 Elements of Dec-POMDPs:

1. **Agents:** Dec-POMDPs involve multiple agents denoted by index  $i \in \{1, 2, \dots, N\}$ , each with its own local observation space  $O_i$  and action set  $A_i$ . The agents interact with the environment and communicate with each other to make coordinated decisions.
2. **States:** The environment's underlying state is not directly observable. Instead, each agent  $i$  receives partial, private observations  $o_i$  that provide probabilistic information about the true state. The state space is denoted as  $S$ .
3. **Observations:** Each agent's observation  $o_i$  is a partial and noisy representation of the global state. Observations are generated based on the current state  $s$  and the agent's local observation function denoted as  $O_i(s)$ .
4. **Actions:** Agents take actions  $a_i$  based on their local observations and communication with other agents. The joint action of all agents affects the evolution of the underlying state. The joint action of all agents is denoted as  $\mathbf{a} = (a_1, a_2, \dots, a_N)$ .
5. **Transition Model:** The transition model  $P(s'|s, \mathbf{a})$  describes the probability distribution over next states  $s'$  given the current state  $s$  and joint actions of all agents  $\mathbf{a}$ .
6. **Observation Model:** The observation model  $P(o_i|s)$  describes the probability distribution over observations  $o_i$  given the current state  $s$  for agent  $i$ .
7. **Global Objective Function:** The goal of the agents is to optimize a global objective function  $J(\mathbf{a})$  that captures the overall mission or task. This objective function depends on the collective actions of all agents and may involve cooperation and coordination.

### 3.2 Gaussian Distribution

The Gaussian distribution Fig. 3.1, also known as the normal distribution, is a continuous probability distribution widely used to model random variables in various real-world scenarios. It is characterized by its mean  $\mu$  and variance  $\sigma^2$ . The probability density function (PDF) of a Gaussian distribution is given by:

$$f(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}, \quad (3.1)$$

where  $x$  is a random variable,  $\mu$  is the mean (expected value), and  $\sigma^2$  is the variance.

The Gaussian distribution is symmetric around its mean, and its shape is determined by the variance. When the variance is small, the distribution is more concentrated around the mean, while a larger variance leads to a wider distribution.




Figure 3.1: Normal Distribution

### 3.3 Policy Gradient Method with Gaussian Approximation:

In RL, the policy gradient method is a popular approach to learning the policy of an agent in a continuous action space. The policy gradient method directly optimizes the policy parameters to maximize the expected cumulative reward.

### 3.3.1 Parameterized Policy

To handle continuous actions, the policy is often parameterized as a Gaussian distribution with a mean vector  $\boldsymbol{\mu}$  and a variance vector  $\boldsymbol{\sigma}^2$ , or alternatively, using the standard deviation  $\boldsymbol{\sigma}$ . The policy function is defined as:

$$\pi_{\boldsymbol{\theta}}(a|s) = \mathcal{N}(\boldsymbol{\mu}(s; \boldsymbol{\theta}), \text{diag}(\boldsymbol{\sigma}(s; \boldsymbol{\theta}))), \quad (3.2)$$

where  $\pi_{\boldsymbol{\theta}}(a|s)$  is the policy function with parameters  $\boldsymbol{\theta}$ , and  $\text{diag}(\cdot)$  extracts the diagonal elements to form a diagonal covariance matrix.

### 3.3.2 Prediction of Mean Using Neural Network

In practice, the mean  $\boldsymbol{\mu}(s; \boldsymbol{\theta})$  is often predicted using a neural network with state  $s$  as input and the policy parameters  $\boldsymbol{\theta}$  as weights. The neural network learns to map states to action means, allowing the agent to approximate actions based on observations from the environment.

### 3.3.3 Adjusting Standard Deviation Over Time

During the training process, the standard deviation  $\boldsymbol{\sigma}(s; \boldsymbol{\theta})$  can be adjusted over time. It is often initialized with some initial value and then decayed towards a minimum value [5], [33], [42]. This adjustment allows the agent to explore more at the beginning of training (high uncertainty) and become more confident in its actions as training progresses (low uncertainty).

## 3.4 Beta Distribution

The Beta distribution Fig. 3.2 is a continuous probability distribution defined on the interval  $[0, 1]$ . It is commonly used to model random variables that represent probabilities or proportions. The Beta distribution is parameterized by two shape parameters, denoted as  $\alpha$  and  $\beta$ , which control the shape of the distribution. The PDF of the Beta distribution is given by:

$$f(x|\alpha, \beta) = \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1}, \quad (3.3)$$

where  $x$  is a random variable,  $\alpha$  and  $\beta$  are the shape parameters, and  $B(\alpha, \beta)$  is the Beta function:

$$B(\alpha, \beta) = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha + \beta)}, \quad (3.4)$$

where  $\Gamma(\cdot)$  is the gamma function.

The Beta distribution is asymmetric and its shape is determined by the values of  $\alpha$  and  $\beta$ . When both  $\alpha$  and  $\beta$  are greater than 1, the distribution is unimodal and skewed. The mode of the distribution is at  $\frac{\alpha-1}{\alpha+\beta-2}$ .




Figure 3.2: Beta Distribution

### 3.5 Policy Gradient Method with Beta Distribution:

In the policy gradient method for RL, the policy is often parameterized as a Beta distribution to handle continuous actions within the range  $[0, 1]$ .

#### 3.5.1 Predicting Alpha and Beta Values from Neural Network:

To approximate the Beta distribution for policy generation, a neural network is used to predict the  $\alpha$  and  $\beta$  parameters. The neural network takes the state  $s$  as input and

produces the  $\alpha$  and  $\beta$  values as output. These values are then used to define the Beta distribution for sampling actions.

### 3.5.2 Interpretation of $1 - \alpha$ and $1 - \beta$ :

The  $\alpha$  and  $\beta$  values can be interpreted as counts of successes and failures, respectively. Specifically, if  $\alpha > 1$  and  $\beta > 1$ , then  $\alpha - 1$  and  $\beta - 1$  can be thought of as the number of successes and failures, respectively. This interpretation is useful when modeling probabilities or proportions in the context of the policy gradient method.

## 3.6 Beta Policy

We propose a straightforward neural network architecture, specifically a Multilayer Perceptron (MLP), with 2 hidden layers each having a size of 256 neurons as we can see in Fig. 3.3. This neural network is then divided into two separate channels, each having a hidden layer with 256 neurons. These channels are designed to predict values for steering angle and acceleration, respectively, in order to control the vehicle.




Figure 3.3: Beta Policy

The implementation of the neural network is carried out using PyTorch [43], a popular deep learning library. The input to the network Fig. 3.3 consists of ego states, which include relevant information about the vehicle in relation to its environment, such as

velocity, checkpoint positions, time, and distance to the destination. Additionally, LiDAR data from the left, right, and front sides of the vehicle is collected and fed into the network. LiDAR is a sensing technology used in many AVs, and its utilization allows the proposed algorithm to be easily adapted and applied in real-world scenarios. The final layer that output these values uses softplus activation function and 1 is added to make sure that the output is always positive. The final layer is  $f(x) = 1 + \log(1 + e^x)$ . As the environment accepts action with  $a = [sa, acc]^T = [-1, 1]^2$ , the output from Beta's PDF is then mapped with the following equation:

$$sa = 2h(\beta_{sa}, \alpha_{sa}) - 1 \quad (3.5)$$

$$acc = 2h(\beta_{acc}, \alpha_{acc}) - 1 \quad (3.6)$$

where  $h(\beta, \alpha)$  is the mean of the Beta distribution with parameters  $\beta$  and  $\alpha$ ,  $sa$  stands for steering angle, and  $acc$  stands for acceleration. It should be noted that we only consider the case where  $\alpha > 1$  and  $\beta > 1$ .

### 3.7 Advantage Function (Experience Sharing)

The advantage function is a key component of the policy gradient method.




Figure 3.4:  $Q$  value Approximator (Critics)

It is used to estimate the advantage of taking an action  $a$  in state  $s$  over taking an action sampled from the policy  $\pi$  in state  $s$ . The advantage function is defined as:

$$A^{\pi_{\theta_k}}(s_t, s_{t+1}, A_j, A_{j+1}, R_G) = R_G + Q(s_{t+1}, A_{j+1}) - Q(s_t, A_j) \quad (3.7)$$

where  $R_G$  is the global reward defined as  $R_G = \frac{1}{T} \sum_{t=0}^T R_k$  whose  $k$  is the individual reward calculated by Monte Carlo Method [33],  $T$  is the total number of agents.  $Q(\star)$  is



the critics part which is popularized by Asynchronous Advantage Actor Critic (A3C) [6] and also adopted by PPO [7],  $Q(s_t, A_j)$  is the Q-value of taking action  $A_j$  in state  $s_t$ , and  $A_{j+1}$  is the action sampled from the policy  $\pi$  in state  $s_{t+1}$ . A 3 layers multi perception layers (MLP) with 256 neurons each is used to approximate the  $Q$  value as we can see in Fig. 3.4.

### 3.8 Objective Function

The objective function is used to optimize the policy parameters  $\theta$ . The goal of the agent is to maximize the culminative reward over time by updating the parameterized policy with:

$$\theta_{k+1} = \arg \max_{\theta} \frac{1}{|D_k|T} \sum_{\tau \in D_k} \sum_{t=1}^T \min\left(\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)} A^{\pi_{\theta_k}}, g(\epsilon, A^{\pi_{\theta_k}})\right) \quad (3.8)$$

where  $D_k$  is the set of trajectories collected by the agent at iteration  $k$ ,  $T$  is the length of the trajectory,  $\pi_{\theta}$  is the policy parameterized by  $\theta$ , and  $g(\epsilon, A^{\pi_{\theta_k}})$  is the clipped surrogate objective function defined as:

$$g(\epsilon, A^{\pi_{\theta_k}}) = \begin{cases} (1 + \epsilon)A^{\pi_{\theta_k}} & \text{if } A^{\pi_{\theta_k}} \geq 0 \\ (1 - \epsilon)A^{\pi_{\theta_k}} & \text{otherwise} \end{cases} \quad (3.9)$$

where  $\epsilon$  is a hyperparameter that controls the magnitude of the update.

### 3.9 Centralized Training and Decentralized Execution

CTDE Fig. 3.5 starts with several agents being placed in the environment. The agents are allowed to run for several predefined steps (2000 timestamps in our case). In the process, samples are collected in a rollout buffer, the buffer then performs some preprocessing tasks such as arranging data and splitting data into experience batches. Those batches are then fed into the central agent's actor-critics network to update the policy by calculating loss and through back propagation. The updated policy is then distributed back to the individual agent. The cycle repeats until the reward feedback converges.

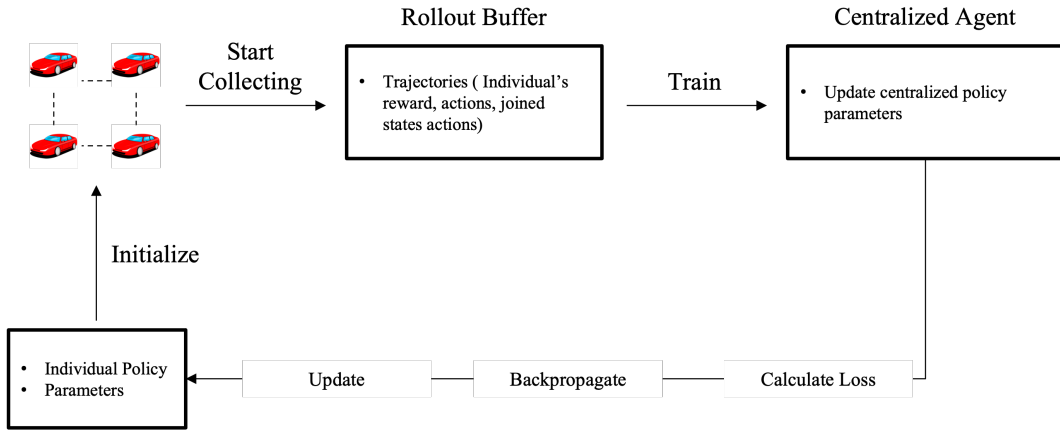


Figure 3.5: Centralized Training and Decentralized Execution

## Chapter 4

# Result and Analysis

### 4.1 Experimental Setup

The following table Table 4.1 describe computer specifications used in the experiment.

Table 4.1: Experiment Setup with PC Specifications

PC Specifications	Details
Processor	Intel Core i7 8700K 6 cores / 12 threads
RAM	32GB
Graphics Card	NVIDIA GeForce GTX 1070 8GB VRAM
Operating System	Ubuntu 22.10 Linux 5.13
Python Version	Python 3.7 Anaconda distribution

In this section, we conducted experiments that compare our PPO Beta and standard PPO Gaussian implementation. We used MetaDrive [41] for traffic simulator Fig. 2.3 which support both single agent and multi-agent training with integrated RL interface (Gym) [44]. We proposed 3 evaluation metrics such as training rewards, safety rate, and success rate. Training reward is the feedback given by the simulator as the score. When the agent drives and reach the destination, 130 score (total reward) is given but we can

set the solving threshold as 80 as average solving score as it implies that the agent already passes through the intersection safely. For safety and success rate, they are collected in evaluation stage which we loaded the pretrained model into the agent as they are in real situation and let the agent run without updating the policy. The success is the ratio of how many times the agent reaches the destination over the number of agents participating in the evaluation. Safety is the ratio of the agents successfully driving on the road without crashes to other objects or vehicles. We run the training for both algorithms for 5 tryouts and plot the average for each evaluation metric. The PPO hyperparameters are shown in table 4.2.

Table 4.2: PPO Configuration

<b>PPO Method</b>	<b>Beta (Ours)</b>	<b>Gaussian</b>
Clipping	0.02	0.02
Gamma	0.99	0.99
Action decay rate	N/A	0.005
Action initial Std	N/A	0.6
Minimum decay	N/A	0.1
K Epoch	5	5
Actor learning rate	0.0003	0.0003
Critic learning rate	0.001	0.001

Here is the explanation of the hyperparameters: Clipping is the clipping value for the PPO loss function. Gamma is the discount factor for the reward. Action decay rate is the decay rate for the action standard deviation. Action initial Std is the initial standard deviation for the action. Minimum decay is the minimum decay rate for the action standard deviation. K Epoch is the number of epochs for the PPO update. Actor learning rate is the learning rate for the actor network. Critic learning rate is the learning rate for the critic network.

## 4.2 Reward Shaping

We designed the reward function  $R$  as:

$$R = c_1(d_t - d_{t-1}) + c_2 \frac{v_t}{v_{max}} + R_T \quad (4.1)$$

where  $d_t$  is the distance between the agent and the destination at time  $t$ ,  $v_t$  is the speed of the agent at time  $t$ ,  $v_{max}$  is the maximum speed of the agent,  $R_T$  is the termination reward, and  $c_1$  and  $c_2$  are the coefficients for the distance and speed reward respectively. We set  $c_1 = 1$  and  $c_2 = 1$ . The design is to encourage the agent to drive faster and reach the destination as soon as possible.  $R_T$  is the termination reward which can be one of the following:

- $R_T = 10$  if the agent reaches the destination.
- $R_T = -5$  if the agent drives out of the road.
- $R_T = -5$  if the agent collides with other vehicles.
- $R_T = -5$  if the agent collides with other objects.

### 4.3 Single Agent Result

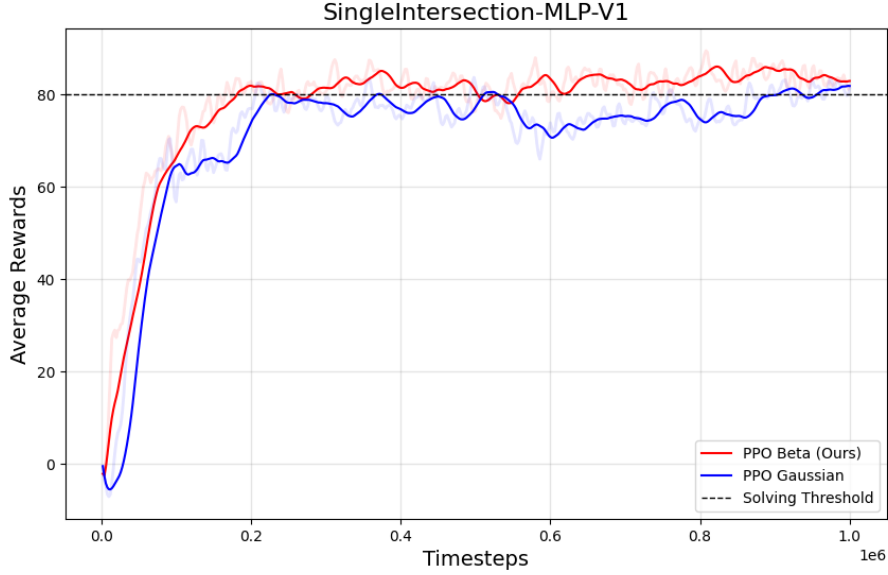


Figure 4.1: Comparison of single agent training

Comparison by average reward between PPO Beta (proposed method) and standard PPO Gaussian. The training is run for 1 million steps. We can see that PPO Beta has higher reward than standard PPO Gaussian which reach the solving threshold (80) faster. The training performance is also more stable than standard PPO Gaussian.

The Gaussian distribution implemented in standard PPO has an adjustment nature that can lead to a drop in training performance after the update of the value function as we can see in Fig. 4.1. This is because the Gaussian distribution is centered around the mean, and when the mean is updated, the distribution must shift to accommodate the new mean. This shift can cause a temporary drop in performance as the agent adjusts to the new distribution. The Beta distribution, on the other hand, does not suffer from this problem. The Beta distribution is more flexible than the Gaussian distribution, and it can accommodate changes in the mean without causing a significant drop in performance. This is because the Beta distribution is not centered around a single point, but rather it is spread out over a range of values. This makes it more robust to changes in the mean.




Figure 4.2: Comparison of single agent training by average success rate

We ran the evaluation for 100 episodes. We can see that PPO Beta has higher success rate than standard PPO Gaussian.

In our experiments, we found that the Beta policy learned to reach the solving threshold (80) faster than the Gaussian policy. Additionally, the average score of the PPO Beta agent was slightly better than the PPO Gaussian agent. These results suggest that the Beta distribution is a more effective choice for continuous control tasks than the Gaussian distribution.




Figure 4.3: Comparison of single agent training by average safety rate

We ran the evaluation for 100 episodes. We can see that PPO Beta has higher safety rate than standard PPO Gaussian.

We evaluated the performance of PPO Beta and PPO Gaussian in a single-agent environment where the training agent had difficulty predicting the actions of the other agents, which were controlled by non-player characters (NPCs). We found that PPO Beta outperformed PPO Gaussian on both the success Fig. 4.2 and safety rates Fig. 4.3 after running for 100 episodes. As shown in ??, both success rates were below 50%. This is because the single-agent environment is inherently difficult, as the training agent cannot coordinate its actions with the other agents. Despite the difficulty of the environment, PPO Beta was able to learn to dodge the incoming vehicle better than PPO Gaussian. This is evident in the safety rate, which was 10% higher for PPO Beta than for PPO Gaussian. These results suggest that PPO Beta is a more effective choice for single-agent environments where the training agent has difficulty predicting the actions of other agents.

## 4.4 Multi Agent Result




Figure 4.4: Comparison of multi-agent training by average reward

Comparison by average reward of 4 agents between PPO Beta (proposed method) and standard PPO Gaussian. The training is run for 700k steps. PPO with Gaussian distribution struggles to learn the environment and the reward is unstable.



We train the multi-agent environment with 4 agents. Here is the interpretation of the result in Fig. 4.4.

#### 4.4.1 Multi-agent PPO Beta

- **Average Reward:** The average reward obtained by the agents over all episodes is approximately 120. This indicates that, on average, the agents perform reasonably well in the environment when using the PPO beta approach.
- **Reward Fluctuations:** The rewards fluctuate throughout the training process. This fluctuation might be caused by various factors, such as the complexity of the environment, the number of agents interacting, and the training dynamics of PPO beta. It is normal to observe some variance in rewards during training.
- **Learning Stability:** The training curve does not show any significant spikes or crashes, which suggests that the training process using PPO beta is stable.

#### 4.4.2 Multi-agent PPO with Gaussian Distribution

- **Average Reward:** The average reward is lower compared to the multi-agent PPO beta. It indicates that using the Gaussian distribution approach might not be as effective in achieving high rewards as the PPO beta in this specific environment.
- **Reward Fluctuations:** Similar to PPO beta, the rewards fluctuate, but it seems to have higher variance with this approach. It is possible that the Gaussian distribution has a higher sensitivity to certain parameters or initial conditions.
- **Learning Stability:** As with PPO beta, there are no apparent stability issues with this approach.

Overall, from this data, it seems that the multi-agent PPO beta performs better in this particular environment compared to the multi-agent PPO with Gaussian distribution. It is worth noting that RL experiments can be sensitive to hyperparameters and initial conditions, so it is essential to perform multiple runs with different seeds to assess the robustness of the results. In this case, we run each experiment with 5 tryouts and plot the average reward over all runs. This helps to reduce the variance in the results and provides a more accurate representation of the performance of each approach.

## Chapter 5

# Conclusion

We proposed multi-agent PPO Beta with dual channel architecture to predict both alpha and beta values for environment with 2 bounded continuous action space. We chose PPO as RL algorithm for training the agent. After testing in traffic intersection environment, the agent can learn to be more careful when crossing the intersection than agent trained with standard Gaussian implementation. We also see that the training performance hardly drops in our implementation. In the future, we are going to extend to multi-agent version with more complex agents' interaction and testing with real street data capture from drone footage such as the data provided by German transportation companies [45] to see how the agents deal with real situations.

# References

- [1] J. Chen, S. E. Li, and M. Tomizuka, “Interpretable end-to-end urban autonomous driving with latent deep reinforcement learning,” 1 2020. [Online]. Available: <http://arxiv.org/abs/2001.08726>
- [2] M. Toromanoff, E. Wirbel, and F. Moutarde, “End-to-end model-free reinforcement learning for urban driving using implicit affordances,” 2019.
- [3] H. Xu, Y. Gao, F. Yu, and T. Darrell, “End-to-end learning of driving models from large-scale video datasets,” 2016.
- [4] B. Peng, Q. Sun, S. E. Li, D. Kum, Y. Yin, J. Wei, and T. Gu, “End-to-end autonomous driving through dueling double deep q-network,” *Automotive Innovation*, vol. 4, pp. 328–337, 8 2021. [Online]. Available: <https://link.springer.com/10.1007/s42154-021-00151-3>
- [5] A. P. Capasso, P. Maramotti, A. Dell ’Eva, and A. Broggi, “End-to-end intersection handling using multi-agent deep reinforcement learning,” 2021.
- [6] V. Mnih and et al., “Asynchronous methods for deep reinforcement learning,” 2016.
- [7] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” 2017.
- [8] J. Schulman, S. Levine, P. Moritz, M. I. Jordan, and P. Abbeel, “Trust region policy optimization,” 2015.
- [9] K. Zhang, Z. Yang, and T. Başar, “Multi-agent reinforcement learning: A selective overview of theories and algorithms,” 2019.

- [10] M. B. Johanson, E. Hughes, F. Timbers, and J. Z. Leibo, “Emergent social learning via multi-agent reinforcement learning,” 2021.
- [11] B. Baker and et al., “Emergent tool use from multi-agent autocurricula,” 2019.
- [12] M. B. Johanson, E. Hughes, F. Timbers, and J. Z. Leibo, “Emergent bartering behaviour in multi-agent reinforcement learning,” 2022.
- [13] S. A. Baby, L. Li, and A. Pokle, “Analysis of emergent behavior in multi agent environments using deep reinforcement learning,” 2022.
- [14] W. Zhou, D. Chen, J. Yan, Z. Li, H. Yin, and W. Ge, “Multi-agent reinforcement learning for cooperative lane changing of connected and autonomous vehicles in mixed traffic,” *Autonomous Intelligent Systems*, vol. 2, no. 1, p. 5, 2022. [Online]. Available: <https://doi.org/10.1007/s43684-022-00023-5>
- [15] S. Wang, H. Fujii, and S. Yoshimura, “Generating merging strategies for connected autonomous vehicles based on spatiotemporal information extraction module and deep reinforcement learning,” *Physica A: Statistical Mechanics and its Applications*, vol. 607, p. 128172, 2022. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0378437122007300>
- [16] K. R. Williams, R. Schlossman, D. Whitten, J. Ingram, S. Musuvathy, J. Pagan, K. A. Williams, S. Green, A. Patel, A. Mazumdar, and J. Parish, “Trajectory planning with deep reinforcement learning in high-level action spaces,” *IEEE Transactions on Aerospace and Electronic Systems*, vol. 59, no. 3, pp. 2513–2529, 2023.
- [17] L. Chen, Z. Jiang, L. Cheng, A. C. Knoll, and M. Zhou, “Deep reinforcement learning based trajectory planning under uncertain constraints,” 2022. [Online]. Available: <https://www.frontiersin.org/articles/10.3389/fnbot.2022.883562>
- [18] M. Fintz, M. Osadchy, and U. Hertz, “Using deep learning to predict human decisions and using cognitive models to explain deep learning models,” *Scientific Reports*, vol. 12, no. 1, p. 4736, 2022. [Online]. Available: <https://doi.org/10.1038/s41598-022-08863-0>
- [19] X. Qian, F. Altché, J. Grégoire, and A. de La Fortelle, “Autonomous intersection management systems: criteria, implementation and evaluation,” *IET Intelligent*

- Transport Systems*, vol. 11, no. 3, pp. 182–189, 2017. [Online]. Available: <https://ietresearch.onlinelibrary.wiley.com/doi/abs/10.1049/iet-its.2016.0043>
- [20] Z. Peng, Q. Li, K. M. Hui, C. Liu, and B. Zhou, “Learning to simulate self-driven particles system with coordinated policy optimization,” 2021.
- [21] T. Wu, M. Jiang, and L. Zhang, “Cooperative multiagent deep deterministic policy gradient (comaddpg) for intelligent connected transportation with unsignalized intersection,” 2020.
- [22] W. He and Y. Huang, “Real-time energy optimization of hybrid electric vehicle in connected environment based on deep reinforcement learning,” *IFAC-PapersOnLine*, vol. 54, no. 10, pp. 176–181, 2021, 6th IFAC Conference on Engine Powertrain Control, Simulation and Modeling E-COSM 2021. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S2405896321015627>
- [23] H. Park and Y. Lim, “Reinforcement learning for energy optimization with 5g communications in vehicular social networks,” 2020. [Online]. Available: <https://www.mdpi.com/1424-8220/20/8/2361>
- [24] E. Walraven, M. T. Spaan, and B. Bakker, “Traffic flow optimization: A reinforcement learning approach,” *Engineering Applications of Artificial Intelligence*, vol. 52, pp. 203–212, 2016. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0952197616000038>
- [25] S. Shahi and H. Lee, “Autonomous rear parking via rapidly exploring random-tree-based reinforcement learning,” 2022. [Online]. Available: <https://www.mdpi.com/1424-8220/22/17/6655>
- [26] S. Yang and B. Yang, “An inductive heterogeneous graph attention-based multi-agent deep graph infomax algorithm for adaptive traffic signal control,” *Information Fusion*, vol. 88, pp. 249–262, 12 2022.
- [27] T. Wang, J. Cao, and A. Hussain, “Adaptive traffic signal control for large-scale scenario with cooperative group-based multi-agent reinforcement learning,” *Transportation Research Part C: Emerging Technologies*, vol. 125, p.

103046, 2021. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0968090X21000760>

- [28] M. Essa and T. Sayed, “Self-learning adaptive traffic signal control for real-time safety optimization,” *Accident Analysis & Prevention*, vol. 146, p. 105713, 2020. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0001457520305388>
- [29] D. Isele, R. Rahimi, A. Cosgun, K. Subramanian, and K. Fujimura, “Navigating occluded intersections with autonomous vehicles using deep reinforcement learning,” in *2018 IEEE International Conference on Robotics and Automation (ICRA)*, 2018, pp. 2034–2039.
- [30] S. Li, K. Shu, C. Chen, and D. Cao, “Planning and decision-making for connected autonomous vehicles at road intersections: A review,” *Chinese Journal of Mechanical Engineering*, vol. 34, p. 133, 2021.
- [31] U. Gunarathna, S. Karunasekara, R. Borovica-Gajic, and E. Tanin, “Intelligent autonomous intersection management,” 2022.
- [32] W.-L. Chen, K.-H. Lee, and P.-A. Hsiung, “Intersection crossing for autonomous vehicles based on deep reinforcement learning,” in *2019 IEEE International Conference on Consumer Electronics - Taiwan (ICCE-TW)*, 2019, pp. 1–2.
- [33] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*. MIT Press, 2018.
- [34] V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, D. Wierstra, and M. Riedmiller, “Playing atari with deep reinforcement learning,” 2013.
- [35] Y. Yang, R. Luo, M. Li, M. Zhou, W. Zhang, and J. Wang, “Mean field multi-agent reinforcement learning,” 2018.
- [36] Y. C. Tang, “Towards learning multi-agent negotiations via self-play,” 2020.
- [37] A. OroojlooyJadid and D. Hajinezhad, “A review of cooperative multi-agent deep reinforcement learning,” 2019.
- [38] J. Siekmann and W. Wahlster, *Agent and Multi-Agent Systems: Technologies and Applications*. Springer, 1998.

- [39] M. Zhou and et al., “Smarts: Scalable multi-agent reinforcement learning training school for autonomous driving,” 2020.
- [40] E. Leurent, “An environment for autonomous driving decision-making,” <https://github.com/eleurent/highway-env>, 2018.
- [41] Q. Li, Z. Peng, L. Feng, Q. Zhang, Z. Xue, and B. Zhou, “Metadrive: Composing diverse driving scenarios for generalizable reinforcement learning,” 2022.
- [42] V. Mnih and et al., “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.
- [43] A. Paszke and et al., “Pytorch: An imperative style, high-performance deep learning library,” 2019.
- [44] G. Brockman and et al., “Openai gym,” 2016.
- [45] J. Bock, R. Krajewski, T. Moers, S. Runde, L. Vater, and L. Eckstein, “The ind dataset: A drone dataset of naturalistic road user trajectories at german intersections,” 2019.